Command	Description	Key word
kubectl get nodes kubectl get nodes -o wide	To see the workers running in the cluster. With -o wide I can see more information. (Internal IP, External IP and more.)	See / Get / Nodes
kubectl get namespaces kubectl get ns	To see the namespaces (You can divide the cluster in multiple namespaces)	See / Get / Namespace
kubectl create namespace nameOfTheNamespace	To create a namespace. Also, I can create a yaml file and run the command: kubectl apply -f namespace.yml namespace.yml `apiVersion: v1` `kind: Namespace` `metadata:` `  name: curso-namespace` Other example: `kind: Namespace` `apiVersion: v1` `metadata:` `  name: testing`	Create / Namespace
kubectl delete namespace nameOfTheNamespace	To delete a namespace. Also, I can delete with a yaml file and run the command: kubectl delete -f namespace.yml	Delete / Namespace
kubectl -n nameOfTheNamespace apply -f pod.yml Note: You run a script always in the same way.	To create a pod in a namespace: pod.yml `apiVersion: v1` `kind: Pod` `metadata:` `  name: wildfly` `spec:` `  containers:` `  - name: wildfly` `  image: jboss/wildfly`	Create / Pod
kubectl get pods	To see the pods	See / Get / Pod
kubectl label pod nameOfThePod role=newRole --overwrite	To change the role in the pod.	Change
kubectl describe pod nameOfThePod	To see information about the pod	See / Describe / Pod
kubectl exec -it nameOfThePod -- /bin/bash	To access to the pod and there you can run commands.	Access / Exec / Pod
kubectl logs -f nameOfThePod	To check the logs of our pod	Logs / Pod

<table>
<tr>
<td>kubectl -n nameOfTheNamespace apply -f deployment.yml

Note: You run a script always in the same way.</td>
<td>To deploy you need to run the command and the yml is:

You create pods with the deployment.

deployment.yml
<pre>apiVersion: apps/v1
kind: Deployment
metadata:
  name: wildfly-deployment
spec:
  selector:
    matchLabels:
      app: wildfly
  replicas: 1
  template:
    metadata:
      labels:
        app: wildfly
    spec:
      containers:
      - name: wildfly
        image: jboss/wildfly:15.0.0.Final
        ports:
        - containerPort: 8080</pre>
deploymentMariadb.yml
<pre>apiVersion: apps/v1
kind: Deployment
metadata:
  name: mariadb-deployment
spec:
  selector:
    matchLabels:
      app: mariadb
  replicas: 1
  template:
    metadata:
      labels:
        app: mariadb
    spec:
      containers:
      - name: mariadb
        image: mariadb
        ports:
        - containerPort: 3306
        env:
          - name: MYSQL_ROOT_PASSWORD
            value: "123"</pre>
You can see environments variables in the mariadb yaml

Other example:</td>
<td>Deploy /
Repplication Controller</td>
</tr>
<tr>
<td></td>
<td><pre>apiVersion: v1
kind: ReplicationController
metadata:
  name: wordpress
spec:
  replicas: 1
  template:
    metadata:
      labels:
        role: wordpress
    spec:
      containers:
      - name: wordpress
        image: wordpress:php7.1-apache
        imagePullPolicy: IfNotPresent
        ports:
        - containerPort: 80</pre></td>
<td></td>
</tr>
</table>

<table>
<tr><td></td><td>Check that kind is different, check the kubernete documentation to know the differentes controllers, but with deployment you change something and create a new version and all the pods in the old version will be replaced with new pods with the new version

<pre>apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: hello
spec:
  replicas: 3
  template:
    metadata:
      labels:
        role: hello
    spec:
      containers:
      - name: hello
        image: gcr.io/google-samples/hello-app:1.0
        imagePullPolicy: IfNotPresent
        ports:
        - containerPort: 8080</pre></td><td></td></tr>
<tr><td></td><td>Also you can assign CPU and Memory to each pod:

<pre>spec:
  containers:
  - name: hello
    image: gcr.io/google-samples/hello-app:1.0
    imagePullPolicy: IfNotPresent
    ports:
    - containerPort: 8080
    env:
    - name: MYSQL_ROOT_PASSWORD
      value: "password"
    resources:
      requests:
        memory: "64Mi"
        cpu: "200m"
      limits:
        memory: "128Mi"
        cpu: "500m"</pre></td><td></td></tr>
<tr><td>kubectl get deployments</td><td>To see the deployments</td><td>See / Deploy</td></tr>
<tr><td>kubectl delete -f deployment.yml</td><td>To delete a deployment.</td><td>Delete / Deploy</td></tr>
</table>

kubectl -n nameOfTheNamespace apply -f servicioMariadb.yml Note: You run a script always in the same way.	You can use a service to access to your pods through different protocols (HTTP, TCP). To create a service: servicioMariadb.yml `kind: Service` `apiVersion: v1` `metadata:` `  name: mariadb-service` `spec:` `  selector:` `  app: mariadb` `  ports:` `  - protocol: TCP` `  port: 3306` `  targetPort: 3306` servicioWildfly.yml `kind: Service` `apiVersion: v1` `metadata:` `  name: wildfly-service` `spec:` `  selector:` `  app: wildfly` `  ports:` `  - protocol: TCP` `  port: 8080` `  targetPort: 8080` As you can see, the selector used is the same selector than in the deployment script. In this way I relate the service with the pods created with the deployment script. Other example: NodePort is a way to expose a port, and in this way you create a port in the worker that point to the container that we want. targetPort is the port in the container. nodePort is the port in the node that point to the point 80 in the container. This service will search the pods with the role wordpress and will send the traffc to these pods.	Create / Service

	`` apiVersion: v1 kind: Service metadata: name: wordpress spec: type: NodePort ports: - port: 80 targetPort: 80 nodePort: 30000 selector: role: wordpress ``` Other example: LoadBalancer , with this, the service will conect to your cloud provider and create a load balancer. ``` apiVersion: v1 kind: Service metadata: name: wordpress-lb spec: type: LoadBalancer ports: - protocol: TCP port: 80 targetPort: 80 name: http selector: role: wordpress ```	
kubectl get service kubectl get svc kubectl -n nameOfTheNamespace get svc	To see the services.	See / Service
kubectl describe service nameOfTheService	To see information about the pod	See / Describe / Service

<table>
<tr>
<td>HPA (horizontal pod autoscaling)

Note: You run a script always in the same way.</td>
<td>To control a metric , for example CPU:

Pods with php-apache (check deployment yml) label will be created maintaining the CPU in 50%, min is 1 and max is 10.

<pre>apiVersion: autoscaling/v2beta2
kind: HorizontalPodAutoscaler
metadata:
  name: php-apache
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: php-apache
  minReplicas: 1
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 50</pre></td>
<td>HPA</td>
</tr>
<tr>
<td>Liveness and readiness probes</td>
<td>Liveness probe is for Kubernetes knows if your pod is alive. If the liveness file, kubernetes restart the pod.
Readiness probe is for Kubernetes knows if your pod is ready to receive traffic.

Example 1: Exec, you can exec a command to know if there is a file. In this case we run cat /tmp/healthy, the liveness will fail because we remove the file after 30 seconds.

<pre>apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
  name: liveness-exec
spec:
  containers:
  - name: liveness
    image: k8s.gcr.io/busybox
    args:
    - /bin/sh
    - -c
    - touch /tmp/healthy; sleep 30; rm -rf /tmp/healthy; sleep 600
    livenessProbe:
      exec:
        command:
        - cat
        - /tmp/healthy
      initialDelaySeconds: 5
      periodSeconds: 5</pre></td>
<td>Liveness and readiness probes</td>
</tr>
</table>

	Example 2: HTTP Get, if the file doesn't exist , the http get fail and Kubernetes restart the pod. `apiVersion: v1` `kind: Pod` `metadata:` `  labels:` `  test: liveness` `  name: liveness-http` `spec:` `  containers:` `  - name: liveness` `  image: k8s.gcr.io/liveness` `  args:` `  - /server` `  livenessProbe:` `  httpGet:` `  path: /healthz` `  port: 8080` `  httpHeaders:` `  - name: X-Custom-Header` `  value: Awesome` `  initialDelaySeconds: 3` `  periodSeconds: 3`	
	Exmple 3: Check TCP port, you can check if a port is open `apiVersion: v1` `kind: Pod` `metadata:` `  name: goproxy` `  labels:` `  app: goproxy` `spec:` `  containers:` `  - name: goproxy` `  image: k8s.gcr.io/goproxy:0.1` `  ports:` `  - containerPort: 8080` `  readinessProbe:` `  tcpSocket:` `  port: 8080` `  initialDelaySeconds: 5` `  periodSeconds: 10` `  livenessProbe:` `  tcpSocket:` `  port: 8080` `  initialDelaySeconds: 15` `  periodSeconds: 20`	

Example 4: In this case we run a nginx and check the port , if the port is open the liveness probe will be passed, for the readiness probe is a HTTP request to check if nginx is ready to receive request. If the liveness probe is Ok but the readiness probe is not OK kubernetes won't restart the pod.

```yaml
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  containers:
  - name: nginx
    image: nginx:alpine
    ports:
    - containerPort: 80
    readinessProbe:
      httpGet:
        path: /
        port: 80
      initialDelaySeconds: 5
      periodSeconds: 10
    livenessProbe:
      tcpSocket:
        port: 80
      initialDelaySeconds: 15
      periodSeconds: 20
```

Volumes	You can create many pods (nginx-01, nginx-02, etc) with a yml script and you can see the volume in the host that will be mounted in the pod: I can check the directory in the host and upload a file there. This volumes is a "HostPath Volume" `apiVersion: v1` `kind: Pod` `metadata:` `  name: nginx-01` `  labels:` `  app: nginx` `spec:` `  containers:` `  - image: nginx` `  name: nginx` `  volumeMounts:` `  - mountPath: /usr/share/nginx/html` `  name: www-volume` `  volumes:` `  - name: www-volume` `  hostPath:` `  # directory location on host` `  path: /www` `  # this field is optional` `  type: Directory`	Volumes
	Now I can create a service to access to the pod: `apiVersion: v1` `kind: Service` `metadata:` `  name: nginx` `spec:` `  type: NodePort` `  ports:` `  - port: 80` `  targetPort: 80` `  nodePort: 30000` `  selector:` `  app: nginx`	

With a kubectl get all I can see the pods and service created:

```
NAME           READY   STATUS    RESTARTS   AGE
pod/nginx-01   1/1     Running   0          77s
pod/nginx-02   1/1     Running   0          50s

NAME                 TYPE        CLUSTER-IP       EXTERNAL-IP   PORT(S)        AGE
service/kubernetes   ClusterIP   10.96.0.1        <none>        443/TCP        42m
service/nginx        NodePort    10.101.171.246   <none>        80:30000/TCP   29s
```

Other type of volume is "DownwardAPI" and you can use it to share data of the Kubernete API with the pods through files

```yaml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-02
  labels:
    app: nginx
spec:
  containers:
  - image: nginx
    name: nginx
    volumeMounts:
    - mountPath: /etc/podinfo
      name: podinfo

  volumes:
    - name: podinfo
      downwardAPI:
        items:
          - path: "labels"
            fieldRef:
              fieldPath: metadata.labels
          - path: "annotations"
            fieldRef:
              fieldPath: metadata.annotations
```

And if you enter to the pod created you can check:

```
root@nginx-02:/etc/podinfo# ls
annotations  labels
root@nginx-02:/etc/podinfo# cat labels
app="nginx"root@nginx-02:/etc/podinfo#
```

Other type of volume is "ConfigMap" and you can use to send configuration or files to pods, if you can't send the configuration with an environment variable you can use ConfigMap.

First you apply the configMap and then the pods.

For this case you share the file index.html

This is the configMap:

```yaml
apiVersion: v1
kind: ConfigMap
metadata:
  name: index-html
data:
  index.html: |-
    Hola soy una configmap
```

This is the pod:

```yaml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-02
  labels:
    app: nginx
spec:
  containers:
  - image: nginx
    name: nginx
    volumeMounts:
    - mountPath: /usr/share/nginx/html
      name: index

  volumes:
    - name: index
      configMap:
        name: index-html
        items:
          - key: index.html
            path: index.html
```

Also you can share a config file, for example:

```yaml
apiVersion: v1
kind: ConfigMap
metadata:
  name: logstash-config
  namespace: logging
data:
  logstash.conf: |-
    input {
      http {
        port      => 8080
      }
    }
    filter {
      prune {
        blacklist_values => {
          "log" => "(MYSQL_PASSWORD|AWS_SECRET)"
        }
      }
    }
    output {
      loggly {
        key => "pone-tu-token-de-loggly-aca"
        tag => "logstash,kubernetes"
        host => "logs-01.loggly.com."
        proto => "https"
      }
    }
```

Other type of volume is "PersistentVolumeClaim" is when you want to create a volume in your cloud provider (AWS, Digital ocean, etc):

pvc yaml file:

```yaml
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: nginx-pvc
spec:
  accessModes:
  - ReadWriteOnce
  resources:
    requests:
      storage: 5Gi
  storageClassName: do-block-storage
```

This is a volume of 5 gb with ReadWriteOnce permission and the storageClassName is the library used by Kubernetes to connect to the API of our cloud provider and create the volume (In our case is digital ocean)

	You apply the pvc file and now you use the volume when you create a pod: `apiVersion: v1` `kind: Pod` `metadata:` `  name: nginx-01` `  labels:` `  app: nginx` `spec:` `  containers:` `  - image: nginx` `  name: nginx` `  volumeMounts:` `  - mountPath: /usr/share/nginx/html` `  name: www-volume` `  volumes:` `  - name: www-volume` `  persistentVolumeClaim:` `  claimName: nginx-pvc` You can see that ClaimName is the name of the volume that we created in the pvc yaml I f we enter to the pod we can see the directory lost+found that always this directory is created when is empty: `root@nginx-01:/# cd /usr/share/nginx/html/` `root@nginx-01:/usr/share/nginx/html# ls` `lost+found` `root@nginx-01:/usr/share/nginx/html#`	
	And if you run the command mount you can see that the volume (digital ocean) is mounted : `shm on /dev/shm type tmpfs (rw,nosuid,nodev,noexec,relatime,size=65536k)` `/dev/disk/by-id/scsi-0DO_Volume_pvc-de4bf1be-596c-11e9-8e68-321af75ee3b6 on /usr/share/nginx/html` `tmpfs on /run/secrets/kubernetes.io/serviceaccount type tmpfs (ro,relatime)` `proc on /proc/bus type proc (ro,relatime)` with kubectl get vpc you can see all the volumes , to delete a volume you need to delete before the pods that are using the volume , and then you can delete the volume with the command kubectl delete pvc nameOfThePVC, In our case nginx-pvc	

Environment variable	You have differents ways to manage environment variables, you can hardcode the key and value in the deployment or pod yaml: ```yaml env: - name: MYSQL_ROOT_PASSWORD value: "123" ``` Or you can use a secret yaml, you run the secret yaml in the namespace (you can have a secret in each environment, so you have a different password in each environment): ```yaml apiVersion: v1 kind: Secret metadata: name: misecreto type: Opaque data: password: PasswordSecreta ``` And in the deployment yaml or pod yaml: ```yaml env: - name: ENV2 valueFrom: secretKeyRef: name: misecreto key: password ``` 	Environment variable